


# Infinispan

Vojtěch Juránek

JBoss - a division by Red Hat

11. 10. 2015, ACM RACS, Prague


# Data today



Source: <http://www.couchbase.com/nosql-resources/what-is-no-sql>

- Huge amount of data
  - You can scale up, but sooner or later you'll need to scale out
  - Need for highly scaleable solution also because of cost effectiveness
- Cloud achitecture - everything is ephemeral
- Lots of data is semi-structured or not structured at all (texts, log files, audio, video, click streams . . . )
  - Needed ability to store unstructured data (often for performance reasons)
  - But ability to talk to RDBMS is often needed as well

# Why caching



Source: Part of xkcd #908

# Why in-memory

- Lots of data is needed in real-time (BigData → FastData)
- Some tasks can be completed much faster when data are kept in memory
- Keeping data in memory during processing of whole application stack, not only during processing in one application in the stack
- With data replication you can keep your data only in memory (no need to store them in persistent storage)




- Data grid platform, written in Java
- In-memory, (optionally) schema-less, No-SQL key-value data store
- Distributed cache - offers massive memory
- Elastic and scalable - can run on hundreds of nodes
- Highly available - no SPOF, resilient to node failures
- Concurrent (MVCC)
- Transactional
- Queryable
- Processing for streaming data

- 100% open-source
  - <https://github.com/infinispan>
  - Apache License, v2.0
- Member of couple research projects (funded by the EU)




# Infinispan modes

- Embedded (library, in-VM)



- Client-server (remote)






# Clustering modes


Under the hood leverages JGroups project for clustering.

- Local - no clustering


- Invalidation



- Replicated




- Distributed



# Remote protocols

- Hot Rod
  - hashing and topology aware
  - failover during topology changes
  - smart request routing
- Memcached
- REST



# Hot Rod clients

Compatible with Java and non-Java platforms. Based on Protocol Buffers - Google's data interchange format.

Clients for

- Java
- C#
- C++
- Python
- Ruby

Python and Ruby clients have only basic functionality.

# Cache stores

A way how to store cache content in some external (persistent) storage.

Two modes:

- Synchronous (write-through)
- Asynchronous (write-behind)

Cache stores:

- Single file store and soft-index file store
- JDBC and JPA cache stores
- LevelDB cache store
- Cloud cache store
- Remote store
- ... and others

Also possible to define custom cache store.

# Transactions

- JTA-compliant transactions
- Ensures consistency of data
- Optimistic and pessimistic locking available
- Read committed and repeatable read isolation levels
- Deadlock detection and recovery
- Data versioning

- Needs some data schema (protobuf file or annotations)
- Search for data using data attributes instead of keys item Featured attributes include keyword, range, fuzzy, wildcard, and phrase queries.
- Combine queries and aggregation functions (but doesn't support joins)
- Sort, filter, and paginate query results
- Support for index or non-indexed queries
- Lucene query API or fluent DSL API

- Role based access control
- User authentication
- Node authentication and authorization
- Encryption of communication
- Audit logging
- Integration with LDAP and/or Kerberos server (includes Active Directory)

# Some other features - brief and selective list

- Data eviction
- Full JSR-107 support (Java Temporary Caching API)
- CDI support
- Remote events
- Client near cache
- Rolling upgrades
- Cross data center replication (also Hot Rod clients support failover to another data center)
- Command line interface
- Map-reduce framework and distributed executors



# Recently implemented features: Infinispan 8



- Functional API
- Distributed streaming
- Continuous querying, grouping and aggregation
- New management console
- Integration with Apache Spark and Hadoop
- ... and more

# Demo

## Integration with Apache Spark

# Some projects using Infinispan

- WildFly
- Hibernate
- Apache Camel
- Apache Marmotta
- CapeDwarf
- Immutable
- apiman
- ... and others



<http://infinispan.org/>

**Thank you for your attention!**

**Questions?**