


# Debezium Asynchronous Engine

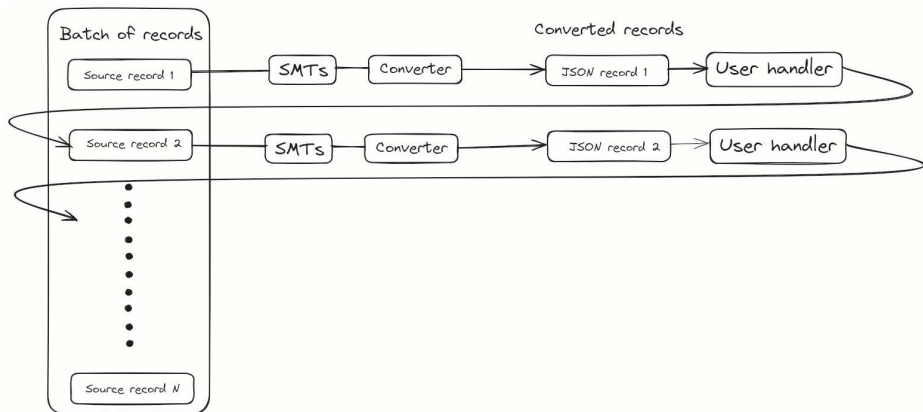
Vojtěch Juránek

Red Hat

May 21st 2024, Debezium F2F meeting, Brno



# Record processing: embedded engine




# Asynchronous engine goals

- Allow to run multiple source tasks for given connector if the connector provides multiple tasks.
- Run potentially time-consuming code (e.g. event transformation or serialization) in the dedicated threads.
- Allow possible further speedup by optionally disabling total ordering of the messages.
- Be well-prepared for future changes and new features:
- Adjust Debezium testsuite to use `DebeziumEngine` interface instead hardcoded `EmbeddedEngine`.

# Asynchronous engine non-goals

- Change `DebeziumEngine` interface.
- Implement any parallelization inside connectors.
- Remove dependency on Kafka Connect API.
- Add support for multiple source connectors or sink connector.

# Debezium asynchronous engine



# Asynchronous engine

- Just another implementation of `DebeziumEngine` interface
- Creation and APIs are same as for `EmbeddedEngine`, only use different builder factory

```
1 DebeziumEngine engine = DebeziumEngine.create(  
2     keyFormat, valueFormat, headerFormat,  
3     ConvertingAsyncEngineBuilderFactory.class.getName  
4     ())  
5     .using(props)  
6     .notifying(consumer)  
7     .build();
```

# Async engine specific configuration options

- `record.processing.threads` - number of threads to be used for processing CDC records.
- `record.processing.shutdown.timeout.ms` - maximum time in milliseconds to wait for processing submitted records when task shutdown is called.
- `record.processing.order` - determines how the records should be produced (ORDERED, UNORDERED).
- `record.processing.with.serial.consumer` - specifies whether the default `ChangeConsumer` should be created from provided `Consumer`, resulting in serial `Consumer` processing.
- `task.management.timeout.ms` - time to wait for task's lifecycle management operations (starting and stopping).




# AsyncEngine internals

# Engine states

AsyncEngine lifecycle phases and states:

- **CREATING** - engine object is being created or was already created, but `run()` method wasn't called yet
- **INITIALIZING** - initializing the connector
- **CREATING\_TASKS** - creating connector tasks
- **STARTING\_TASKS** - starting connector tasks
- **POLLING\_TASKS** - running tasks polling, this is the main phase when the data are produced
- **STOPPING** - the engine is being stopped
- **STOPPED** - engine has been stopped, any call on engine object in this state should fail

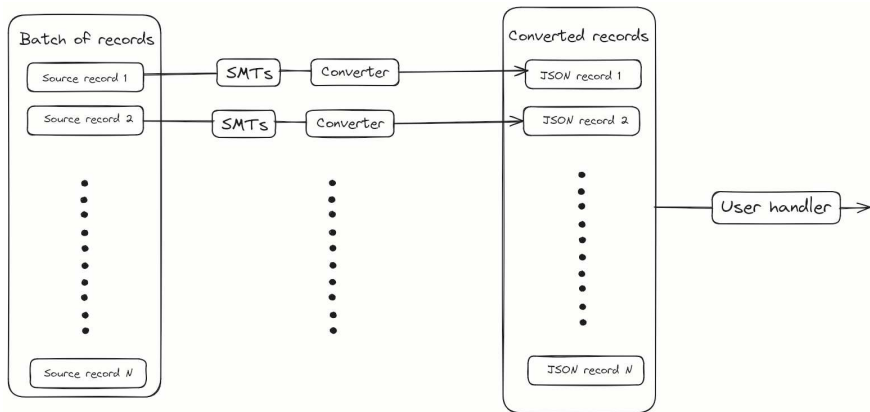
# Debezium asynchronous engine



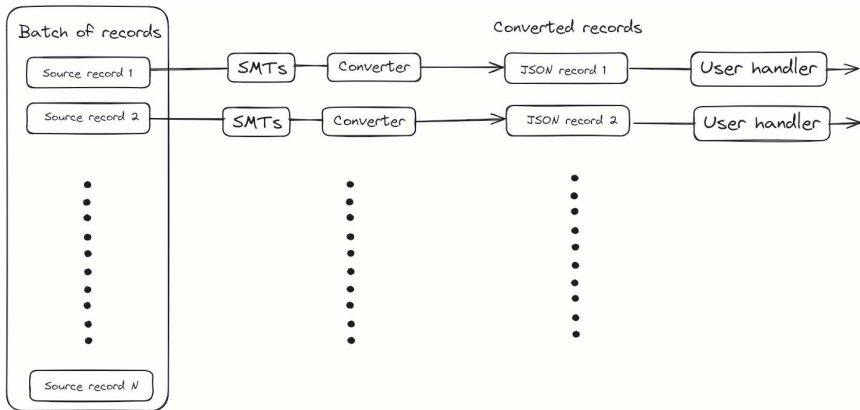
# Record processor

```
1 @Incubating
2 public interface RecordProcessor<R> {
3     void initialize(
4         final ExecutorService recordService,
5         final Transformations transformations,
6         final Function<SourceRecord, R> serializer,
7         final RecordCommitter committer);
8
9     void processRecords(final List<SourceRecord> records)
10        throws InterruptedException;
11 }
```

# Record processing: async. engine ordered



# Record processing: async. engine unordered



# Auxiliary interfaces and classes

```
1 @Incubating
2 public interface DebeziumSourceConnector {
3     DebeziumSourceConnectorContext context();
4     void initialize(DebeziumSourceConnectorContext
5         context);
6 }
```

```
1 @Incubating
2 public interface DebeziumSourceConnectorContext {
3     OffsetBackingStore offsetStore();
4     OffsetStorageReader offsetStorageReader();
5     OffsetStorageWriter offsetStorageWriter();
6 }
```

and more. See

[debezium-embedded/src/main/java/io/debezium/engine/source/](#)

# Code walk through



- gRPC
- virtual threads - switching just few lines of code
- Quarkus integration

- [DDD-7: Asynchronous Debezium Embedded Engine](#)
- [Discussion under DDD-7 PR](#)
- [DBZ-7024](#) - main async. engine tracking Jira
- Other possible interesting Jiras: [DBZ-7764](#), [DBZ-7777](#)

# Thank you!

# Questions?